

NAME: Heet Dhanuka
ROLL NO.: B -34
BATCH: B2

DWDM Practical 4

Aim : Implementing indexes and IOT in Oracle.

Q1. Create a copy of customers table and name it customers_copy_btree_<rollno>.
Create individual b-tree indexes on the following columns of the table
customers_copy_btree_<rollno> :

- (a) cust_gender
- (b) cust_year_of_birth
- (c) cust_last_name
- (d) cust_street_address

```
SQL> CREATE TABLE customers_copy_btree_34 AS  
2  SELECT * FROM customers;  
  
Table created.
```

```
SQL> CREATE INDEX idx_cust_gender_34
  2  ON customers_copy_btree_34(cust_gender);

Index created.

SQL>
SQL> CREATE INDEX idx_cust_year_of_birth_34
  2  ON customers_copy_btree_34(cust_year_of_birth);

Index created.

SQL>
SQL> CREATE INDEX idx_cust_last_name_34
  2  ON customers_copy_btree_34(cust_last_name);

Index created.

SQL>
SQL> CREATE INDEX idx_cust_street_address_34
  2  ON customers_copy_btree_34(cust_street_address);

Index created.
```

Q2. Create bitmap indexes on the above columns. How long does it take to create bitmap indexes? Compare it with the results of btree index creation.

```

SQL> CREATE BITMAP INDEX bm_idx_cust_gender_34
  2  ON customers_copy_btree_34(cust_gender);

Index created.

SQL>
SQL> CREATE BITMAP INDEX bm_idx_cust_year_of_birth_34
  2  ON customers_copy_btree_34(cust_year_of_birth);

Index created.

SQL>
SQL> CREATE BITMAP INDEX bm_idx_cust_last_name_34
  2  ON customers_copy_btree_34(cust_last_name);

Index created.

SQL>
SQL> CREATE BITMAP INDEX bm_idx_cust_street_address_34
  2  ON customers_copy_btree_34(cust_street_address);

Index created.

```

1. **B-Tree Index** is typically **faster** to create than a **Bitmap Index**, especially if the table is large.
2. **Bitmap Index** takes **more time** but is **smaller in size** for columns with low cardinality

Aspect	B-Tree Index	Bitmap Index
Query	CREATE INDEX	CREATE BITMAP INDEX
Best for	High-cardinality columns	Low-cardinality columns
Creation Speed	Faster	Slower
Storage Size	Larger	Smaller
DML Impact	Minimal	Slows updates/deletes
Usage	OLTP (frequent updates)	OLAP (analytics, filters)

Q3. Find the size of each segment: customers_copy_bitmap and customers_copy_btree
(Hint : Use users_segment table)

```
SQL> SELECT segment_name, bytes / (1024 * 1024) AS size_mb
2 FROM user_segments
3 WHERE segment_name IN ('CUSTOMERS_COPY_BTREE_34', 'CUSTOMERS_COPY_BITMAP_34');

SEGMENT_NAME
-----
SIZE_MB
-----
CUSTOMERS_COPY_BTREE_34
.0625
```

Q4. Do as directed :

a. Create function based index on Employee table of HR schema. Function should be on salary attribute based on commission percentage

```
SQL> CREATE INDEX idx_emp_salary_comm ON hr.employees (salary * (1 + NVL(commission_pct, 0)));
Index created.
```

b. Find out list of employees having commission percentage less than 50000.

```
SQL> SELECT * FROM hr.employees WHERE salary * NVL(commission_pct, 0) < 50000;
```

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
100	Steven	King	SKING	515.123.4567	17-JUN-03	AD_PRES	24000

COMMISSION_PCT	MANAGER_ID	DEPARTMENT_ID
90	100	90

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
101	Neena	Kochhar	NKOCHHAR	515.123.4568	21-SEP-05	AD_VP	17000

COMMISSION_PCT	MANAGER_ID	DEPARTMENT_ID
90	100	90

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
102	Lex	De Haan	LDEHAAN	515.123.4569	13-JAN-01	AD_VP	17000

COMMISSION_PCT	MANAGER_ID	DEPARTMENT_ID
60	102	60

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
103	Alexander	Hunold	AHUNOLD	590.423.4567	03-JAN-06	IT_PROG	9000

COMMISSION_PCT	MANAGER_ID	DEPARTMENT_ID
60	103	60

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
104	Bruce	Ernst	BERNST	590.423.4568	21-MAY-07	IT_PROG	6000

COMMISSION_PCT	MANAGER_ID	DEPARTMENT_ID
60	104	60

EMPLOYEE_ID	FIRST_NAME	LAST_NAME	EMAIL	PHONE_NUMBER	HIRE_DATE	JOB_ID	SALARY
105	David	Austin	DAUSTIN	590.423.4569	24-MAY-07	IT_PROG	6000

c. Create function based index on employee name for Upper and lower function.

```
SQL> CREATE INDEX idx_emp_name_upper ON hr.employees (UPPER(last_name));  
Index created.  
  
SQL> CREATE INDEX idx_emp_name_lower ON hr.employees (LOWER(last_name));  
Index created.
```

d. Create user table with attributes (UserId, UserName, Gender)

```
SQL> CREATE TABLE user_table (    UserId NUMBER PRIMARY KEY,    UserName VARCHAR2(100),  
    Gender CHAR(1));  
Table created.
```

e. Insert 10000 records in user table

```
SQL> BEGIN  
2     FOR i IN 1..10000 LOOP  
3         INSERT INTO USERS (UserId, UserName, Gender)  
4             VALUES (i, 'User' || i, CASE WHEN MOD(i,2) = 0 THEN 'M' ELSE 'F' END);  
5     END LOOP;  
6     COMMIT;  
7 END;  
8  
9 /  
  
PL/SQL procedure successfully completed.
```

f. Build regular index on Username

```
SQL> CREATE INDEX idx_user_name  
2   ON USERS(UserName);  
  
Index created.
```

g. Build function based index on user name based on Upper function

```
SQL> CREATE INDEX idx_user_name_upper  
2   ON USERS(UPPER(UserName));  
  
Index created.
```

h. Compare the response time and comment.

```
SQL> SET TIMING ON;  
SQL> SELECT * FROM USERS WHERE UserName = 'User5000';
```

```
      USERID  
-----  
USERNAME  
-----  
G  
-  
      5000  
User5000  
M  
  
Elapsed: 00:00:00.00  
SQL> SET TIMING OFF;  
SQL> |
```

- 1.Regular Index works well when searching exactly as stored.
- 2.Function-Based Index is helpful when using UPPER(), LOWER(), or calculations in WHERE clauses.
- 3.Without Function-Based Index, UPPER(UserName) = 'USER5000' results in a full table scan, slowing down performance.

Q5. Do as directed :

- a. Create an IOT look_ups with the attributes (lookup_code, lookup_value, lookup_description).
- b. Constraint: lookup_code should be primary key
- c. lookup_description should be in overflow area.

```
SQL> CREATE TABLE LOOK_UPS (  
2      LOOKUP_CODE NUMBER PRIMARY KEY,  
3      LOOKUP_VALUE VARCHAR2(100),  
4      LOOKUP_DESCRIPTION VARCHAR2(500)  
5  ) ORGANIZATION INDEX  
6  PCTTHRESHOLD 20  
7  OVERFLOW;
```

Table created.

```
SQL> |
```

Q6. Do as directed :

a. Create a Index Organized Table(IOT) emp_iot based on hr.employees

```
SQL> CREATE TABLE EMP_IOT (  
  2     EMPLOYEE_ID NUMBER PRIMARY KEY,  
  3     FIRST_NAME VARCHAR2(50),  
  4     LAST_NAME VARCHAR2(50),  
  5     EMAIL VARCHAR2(100),  
  6     PHONE_NUMBER VARCHAR2(20),  
  7     JOB_ID VARCHAR2(10),  
  8     SALARY NUMBER,  
  9     COMMISSION_PCT NUMBER,  
 10     MANAGER_ID NUMBER,  
 11     DEPARTMENT_ID NUMBER  
 12 ) ORGANIZATION INDEX;
```

Table created.

b. Create a Index Organized Table(IOT) emp101_emp based on hr.employees. Place the column hiredate in overflow area.


```

SQL> CREATE TABLE EMP101_IOT (
  2     EMPLOYEE_ID NUMBER PRIMARY KEY,
  3     FIRST_NAME VARCHAR2(50),
  4     LAST_NAME VARCHAR2(50),
  5     EMAIL VARCHAR2(100),
  6     PHONE_NUMBER VARCHAR2(20),
  7     JOB_ID VARCHAR2(10),
  8     SALARY NUMBER,
  9     COMMISSION_PCT NUMBER,
10     MANAGER_ID NUMBER,
11     DEPARTMENT_ID NUMBER,
12     HIRE_DATE DATE
13 ) ORGANIZATION INDEX
14 PCTTHRESHOLD 20
15 OVERFLOW;

```

Table created.

c. Compare the timings of executing select all from employees, emp_iot, and emp101_iot. Comment on your observations.

HR.EMPLOYEES (Heap Table) → Slowest, full table scan increases disk I/O.

EMP_IOT (IOT Table) → Faster, rows indexed by primary key for efficient lookups.

EMP101_IOT (IOT with Overflow) → Fastest if HIRE_DATE is rarely accessed; slower if HIRE_DATE is frequently queried due to overflow access.